

# Deep Learning

---

Why Did We Arrive Here (Today)?

# Introduction

---

1. Why such interest?
2. Why now?
3. Most famous example
4. In what conditions is this not overfitting?
5. In what conditions is this overfitting?

DataScience@SMU

# Why Such Interest?

---

# Why Such Interest?

---

- Terminology
  - Neural networks (NN) sound like brains
  - NN reintroduced the terminology of AI
- Disruptive technology
  - Replace human sensing
    - (OCR, voice recognition)
  - Replace human movement
    - (Robots, cars)
- Before, no models were plausible replacements for such tasks.

DataScience@SMU

# Why Now?

---

# Why Now?

---

Since NN are not new, why now

- GPUs
- Aggregation of massive datasets

Secondary

- Stagnating research on other topics
- Let CS engineers contribute



DataScience@SMU

# Most Famous Example

---

# Most Famous Example

---

Convolutional neural networks

- Vectorize an image
- Convolutional NN

Feature representation

- CNN
- GANS
- Autoencoders
- Recurrent neural networks  
LSTM

```
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
from keras import backend as K
batch_size = 128
num_classes = 10
epochs = 12
img_rows, img_cols = 28, 28
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])
model.fit(x_train, y_train,
        batch_size=batch_size,
        epochs=epochs,
        verbose=1,
        validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

DataScience@SMU

# In What Conditions Is This Not Overfitting?

---

# Why Is This Not Overfitting?

---

Accurate collection of distribution

- Claim: Google has a representative sample of all images (i.e., the distribution is stable and there is ample supply).
- Claim: Amazon has representative sample of product reviews (?).
- Claim: Newspapers have a reasonable estimate of how newspapers are written.
- We are using a holdout set that we only look at after we have decided what our optimal model is.
- We have ample enough data to retrain our model whenever we need to, using a new holdout set.
- We've evaluated core techniques for increasing generalizability such as dropout, etc.



DataScience@SMU

# In What Conditions Is This Overfitting?

---

# Why Is This Overfitting?

---

- We never have a stable distribution.
- We never have enough data (too few or too hard to access).
- Engineers often cut corners in model building.
- Companies gain an edge by working at the edge and the stability we require for DL is not there.
- Deep learning memorizes relationships within the distribution but does nothing to generalize this knowledge.
- Malicious attacks are not yet included in the training data!
- When distributions change, how will we know?

DataScience@SMU